

Ridge Regression Lab

Group 2-07

Dataset Hitters (from ISLR)

- ▶ Hitters is a dataset with 263 observations of 20 variables. Here, we take the last variable, salary, as our response and the other 19 variables as predictors.

```
library("ISLR", "glmnet")
Hitters <- na.omit(Hitters)
x <- model.matrix(Salary ~ ., Hitters)[, -1]
y <- Hitters$Salary
```

- ▶ Split the sample by 50/50:

```
set.seed(1)
train <- sample(nrow(x), nrow(x)/2)
test <- -train
y.test <- y[test]
```

glmnet()

- ▶ x & y
- ▶ elasticnet mixing parameter $\frac{1-\alpha}{2}||\beta||_2^2 + \alpha||\beta||_1$ where $\alpha \in [0, 1]$
($\alpha = 0$ for Ridge; $\alpha = 1$ for Lasso)

```
ridge.mod <- glmnet(x,y,alpha=0)
```

- ▶ Note that the *glmnet()* function standardizes the variables so that they are on the same scale. To turn off this default setting, use the argument **standardize=FALSE**.

Choose λ by cross-validation


- ▶ By default the function performs 10-fold cross-validation, though this can be changed using the argument *nfolds*.
- ▶ the *lambda.min* is the lambda that gives $\min(MSE - length(lambda))$

```
set.seed(1)
cv.out <- cv.glmnet(x[train,],y[train],alpha=0)
bestlam <- cv.out$lambda.min
bestlam
```

```
## [1] 211.7416
```

Plot of cv.out

```
plot(cv.out)
```



Final Ridge Regression Model

```
ridge.pred <- predict(ridge.mod,s=bestlam,newx=x[test,])  
predict(ridge.mod,type="coefficients",s=bestlam)[1:20,]
```

##	(Intercept)	AtBat	Hits	HmRun	
##	9.88487157	0.03143991	1.00882875	0.13927624	1.
##	RBI	Walks	Years	CAtBat	
##	0.87318990	1.80410229	0.13074381	0.01113978	0.
##	CHmRun	CRuns	CRBI	CWalks	
##	0.45158546	0.12900049	0.13737712	0.02908572	27.
##	DivisionW	PutOuts	Assists	Errors	Ne
##	-91.63411299	0.19149252	0.04254536	-1.81244470	7.

- Ridge regression does **not** perform variable selection!

Benefit to performing ridge regression with bestlam

```
ridge.train <- glmnet(x[train,],y[train],alpha = 0)
```

- ▶ Remember that when $\lambda = 0$, we are not doing shrinkage regression but least square regression.

```
ridge.pred <- predict(ridge.train,s=0,newx=x[test,],  
                      x=x[train,],y=y[train])  
#mean((ridge.pred-y.test)^2)
```

- ▶ With bestlam

```
ridge.pred <- predict(ridge.train,s=bestlam,newx=x[test,],  
                      x=x[train,],y=y[train])  
#mean((ridge.pred-y.test)^2)
```

cont.


- ▶ In general, if we want to fit a (unpenalized) least squares model, then we should use the `lm()` function, since that function provides more useful outputs, such as standard errors and p-values for the coefficients.

```
coef(lm(y~x, subset=train))
```

```
predict(ridge.mod,s=0,exact=T,type="coefficients",x=x[train,],y=y[train])
```


Comparing different λ and their l2 norm

```
plot(lambda,l2)
```



What's the difference between coefficients?

```
lambda_large<- coef(ridge.mod)[,1]
lambda_small<- coef(ridge.mod)[,50]
cbind(lambda_large,lambda_small)
```

##	lambda_large	lambda_small
## (Intercept)	5.359259e+02	213.066443434
## AtBat	1.221172e-36	0.090095728
## Hits	4.429736e-36	0.371252756
## HmRun	1.784944e-35	1.180126956
## Runs	7.491019e-36	0.596298287
## RBI	7.912870e-36	0.594502390
## Walks	9.312961e-36	0.772525466
## Years	3.808598e-35	2.473494238
## CAtBat	1.048494e-37	0.007597952
## CHits	3.858759e-37	0.029272172
## CHmRun	2.910036e-36	0.217335716
## CRuns	7.741531e-37	0.058705097
## C RBI	7.989430e-37	0.060722036

Another example in classification

- ▶ Setting up

```
yclass <- rep("Yes", length(y))  
yclass[y < median(y)] <- "No"  
yclass <- factor(yclass)  
yclass.test <- yclass[test]
```

- ▶ Use Logistic Regression to compare coefficients
Recall `glm()` uses **family="binomial"**, and call **type = "response"** in `predict()`;
`coef_glm <- logistic.mod.class$coefficients`


Ridge regression (classification)

- Again, override family as binomial in both `glmnet()` and `cv.glmnet()`

```
ridge.mod.class <- glmnet(x, yclass, alpha=0,
                          family="binomial")
cv.out.class <- cv.glmnet(x[train, ], yclass[train],
                         alpha=0, family="binomial")
bestlam <- cv.out.class$lambda.min
ridge.pred.class <- predict(ridge.mod.class, s=bestlam,
                           newx=x[test, ], type="class")
error.rate.ridge <- mean(ridge.pred.class != yclass.test)
ridge.coefficients <- predict(ridge.mod.class,
                              type="coefficients",
                              s=bestlam)[1:20, ]
```

Plot for cv.out.class

```
plot(cv.out.class)
```



Comparison

► Error rate

```
cbind(error.rate.logistic,error.rate.ridge)
```

```
##          error.rate.logistic error.rate.ridge
## [1,]          0.2348485          0.1590909
```

► Coefficients

```
cbind(coef_glm,ridge.coefficients)
```

```
##              coef_glm ridge.coefficients
## (Intercept) -3.551364583      -2.822169e+00
## AtBat        -0.006591710       8.156105e-04
## Hits         0.036790282       3.767020e-03
## HmRun        0.043504900       2.264340e-03
## Runs        -0.064130104       3.922014e-03
## RBI          0.007675413       3.281934e-03
## HmRun        0.054262722       7.247705e-03
```